

LANGMUIR FITTING: QUICK REFERENCE CHEAT SHEET

One-page summary of everything for Phase 2

THE LANGMUIR EQUATION (What We're Fitting)

$$q = (q_{\max} \times K_L \times C_e) / (1 + K_L \times C_e)$$

q = Adsorbed amount (mg/g) ← What we predict
 $q_{\max}$ = Max capacity (mg/g) ← Parameter (changes with conditions)
 K_L = Binding constant (L/mg) ← Parameter (changes with conditions)
 C_e = Equilibrium conc (mg/L) ← Calculated from data

Key insight: $q_{\max}$ and K_L aren't constant—they change with pH, temperature, ions, etc.

THE MODEL WE USE (What's Different)

NOT: Simple 2-parameter Langmuir (ignores factors)

YES: Multi-factor Langmuir with polynomial features

Step 1: Take 10 factors (pH, CO, Time, Dose, Temp, Flow, Cl, Hard, CO3, NOM)

Step 2: Add squares (pH², CO², Time², ...)

Step 3: Add interactions (pH×CO, pH×Time, ...)

Result: 66 features from 10 factors

Step 4: Linear regression on 66 features → 67 coefficients

Result: $q = \text{ + } \times f + \dots + \times f$

Why? Factors change the curve shape → need to capture those changes.

WHAT HAPPENS STEP-BY-STEP

- [1] Load: Read CSV (500 samples, 10 factors)
 ↓
- [2] Standardize: Scale to mean=0, std=1 (numbers become comparable)
 ↓
- [3] Expand: 10 factors → 66 polynomial features
 ↓
- [4] Fit: Linear regression (minimize prediction error)
 ↓
- [5] Predict: Apply model to get predictions for all 500 points
 ↓
- [6] Calculate: R², RMSE, residuals, diagnostics
 ↓
- [7] Visualize: 4 diagnostic plots (save as PNG)
 ↓

[8] Save: CSV with predictions, JSON with metadata

Runtime: ~1 minute (mostly visualization)

EXPECTED RESULTS AT A GLANCE

Metric	Expected	Interpretation
R²	0.84-0.87	Explains 84-87% of variance Good!
RMSE	0.9-1.2 mg/g	Average error ± 0.9 mg/g (13% of range)
MAE	0.7-0.9 mg/g	Typical error magnitude
Residual mean	0	No systematic bias
Residual std	0.92 mg/g	Residual spread (what ML learns)
Residual dist	Normal	Bell-shaped histogram

Bottom line: $R^2 = 0.85$ is EXCELLENT for a chemical model. Leaves 15% for ML to learn.

THE 4 DIAGNOSTIC PLOTS

Panel 1: Actual vs Predicted

Points follow diagonal line?

No systematic curve?

→ If YES: Good fit!

Panel 2: Residuals vs Predicted

Random scatter around zero line?

Constant width (not funnel)?

→ If YES: Homoscedastic, good!

Panel 3: Residual Distribution

Bell curve shape?

Centered at zero?

No extreme outliers?

→ If YES: Normal distribution, good!

Panel 4: Q-Q Plot

Points follow straight line?

Minor deviations at tails OK?

→ If YES: Normal distribution, good!

INTERPRETING YOUR RESULTS

R² Score

What it means: "Model explains X% of variance"

$R^2 = 0.85$ means:

85% of q variation is explained by Langmuir model
15% remains as residuals (for ML to learn)

Is 0.85 good?

YES! For a chemistry model, this is excellent.
Shows Langmuir is appropriate.
Leaves good room for ML improvement (15%).

If $R^2 < 0.75$?

~ Probably OK, but investigate residuals

If $R^2 > 0.95$?

Very good! (But may indicate overfitting in theory)
Actually fine for Phase 2 (synthetic data is cleaner)

RMSE Score

What it means: "Average prediction error is X mg/g"

RMSE = 0.92 means:

Typical prediction off by ± 0.92 mg/g
Range of data is 1.6-8.3 (6.68 total)
Error is $0.92/6.68 = 13.8\%$ of range

Is 0.92 good?

YES! $< 15\%$ of range is good.
Provides practical confidence intervals.

Rule of thumb:

$< 10\%$ of range = Excellent
 $10-20\%$ of range = Good (You are here)
 $20-30\%$ of range = Acceptable
 $> 30\%$ of range = Poor

Residual Analysis

What to look for:

1. Mean 0?

YES $\rightarrow$ No systematic bias
NO $\rightarrow$ Model over/underestimates

2. Standard deviation 0.92?

YES $\rightarrow$ Consistent with RMSE
NO $\rightarrow$ High variability (investigate)

3. Normal distribution?
 - YES → Good for statistics
 - NO → Non-normal (usually OK, regression is robust)
 4. Random scatter (no patterns)?
 - YES → Model captured the pattern
 - NO → Model missing something (e.g., non-linearity)
 5. No extreme outliers?
 - YES → Good data quality
 - NO → Investigate unusual conditions
-

OUTPUT FILES EXPLAINED

langmuir_predictions.csv

Contains: 500 rows, 15 columns

- First 12: Original data (factors + design columns)
- Column 13: q_removal (actual)
- Column 14: q_predicted (model)
- Column 15: residual (actual - predicted)

Use for:

- Phase 3: Analyze residual patterns
- Find worst predictions (largest |residual|)
- Identify where model struggles

langmuir_model_info.json

Contains: Model metadata and metrics

- $R^2 = 0.8456$
- RMSE = 0.9231
- MAE = 0.7145
- residual_mean = -0.000001
- residual_std = 0.9228

Use for:

- Reference baseline
- Compare with Phase 4 ML results
- Track improvement to hybrid model

langmuir_diagnostics.png

Contains: 4-panel diagnostic plot

- Panel 1: Actual vs Predicted
- Panel 2: Residuals vs Predicted
- Panel 3: Residual Distribution
- Panel 4: Q-Q Plot

Use for:

- Visual validation of fit
- Check for systematic patterns
- Identify outliers or issues
- Report figure for documentation

TROUBLESHOOTING QUICK FIXES

Problem	Likely Cause	Quick Fix
$R^2 < 0.70$	Model too simple OR data issue	Check data quality, try degree=3 polynomial
$R^2 > 0.98$	Overfitting (theory)	OK for Phase 2, proceed normally
$RMSE > 2.0$	Model missing major effects	Check residual plots for patterns
Residuals not normal	Outliers or skewness	Check Q-Q plot, identify outliers
Funnel in residual plot	Heteroscedasticity	OK for Phase 2, ML will learn it
ImportError (sklearn)	Package not installed	<code>pip install scikit-learn</code>
File not found	CSV in wrong folder	<code>cp data/dataset_simulated_500.csv .</code>

PHASE 2 SUCCESS CRITERIA

Script runs without errors R^2 is 0.80-0.95 (0.84-0.87 expected) $RMSE$ is < 1.5 mg/g (0.9-1.2 expected) Residuals appear normal (bell curve) No systematic patterns in residuals 3 output files created successfully Diagnostic plots look reasonable You understand what was fitted

If all : You're ready for Phase 3!

WHAT COMES NEXT

Phase 3 (Residual Analysis): - Analyze what Langmuir missed - Identify patterns in residuals
- Engineer features to capture patterns - Prepare data for ML training

Phase 4 (ML Training): - Train Random Forest, XGBoost, MLP - Predict residuals (not absolute values) - Expected: R^2 0.90-0.93 on residuals

Phase 5 (Hybrid Integration): - Combine: $q_{\text{hybrid}} = q_{\text{langmuir}} + q_{\text{ml_correction}}$ - Achieve: R^2 0.94 (Target!) - Success: 20-35% improvement

KEY EQUATIONS TO REMEMBER

Langmuir:

$$q = (q_{\max} \times K_L \times C_e) / (1 + K_L \times C_e)$$

R^2 Score:

$$R^2 = 1 - (SS_{\text{residuals}} / SS_{\text{total}})$$

$$R^2 = 1 - (\Sigma(y - \hat{y})^2 / \Sigma(y - \bar{y})^2)$$

RMSE:

$$RMSE = \sqrt{(\Sigma(y - \hat{y})^2 / n)}$$

Residuals:

$$\text{residuals} = y_{\text{actual}} - y_{\text{predicted}}$$

MENTAL MODEL

Think of Phase 2 like this:

Question: "How much fluoride does the carbon remove?"

Chemistry says:

"Use Langmuir equation with these factors."

We do:

1. Learn how factors modify Langmuir
2. Build a model that captures that
3. Make predictions (should be ~85% accurate)
4. Calculate what the model missed (residuals)

Result:

- Chemistry explains 85% ($R^2 = 0.85$)
- Residuals = 15% (opportunities for ML)

Next:

- ML learns to predict residuals
 - Hybrid model = 85% + 9% = 94% (target!)
-

DECISION TREE: IS MY FIT GOOD?

```
      Is  $R^2$  0.80?  
      |  
YES   |   NO  
      |
```

```

GO TO INVESTIGATE
PHASE 3 (check data)
|
Are residuals
normal?
|
YES | NO
|
GOOD! INVESTIGATE
PHASE 3 (outliers)
|
Are there
patterns in
residuals?
|
YES | NO
|
EXPECTED PERFECT!
(Phase 3 RANDOM
learns)

```

ONE-LINE SUMMARY

Phase 2 Langmuir Fitting: Fit multi-factor Langmuir model to data, achieve $R^2 = 0.85$, understand residuals, prepare for ML.

QUICK STATS TABLE

Parameter	Value	Unit
Samples	500	points
Factors	10	original
Features	66	polynomial
Coefficients	67	(+ 66 's)
Sample/Feature ratio	7.6	(good, avoid overfitting)
Expected R^2	0.84-0.87	fraction
Expected RMSE	0.9-1.2	mg/g
Data range	1.6-8.3	mg/g
Error as % range	13.8%	excellent
Runtime	<1	minute
Output files	3	(CSV, JSON, PNG)
Next phase	Phase 3	Residual Analysis

This is your complete Phase 2 reference guide!

Keep this nearby while running Phase 2. Refer to the comprehensive guides for detailed explanations.

Ready? Run: `python phase2_langmuir_fitting.py`

EOF `cat /mnt/user-data/outputs/LANGMUIR_CHEAT_SHEET.md`